

An ant colony optimization approach to flexible protein–ligand docking

Oliver Korb · Thomas Stützle · Thomas E. Exner

Received: 24 January 2007 / Accepted: 20 August 2007 / Published online: 19 September 2007
© Springer Science + Business Media, LLC 2007

Abstract The prediction of the complex structure of a small ligand with a protein, the so-called protein–ligand docking problem, is a central part of the rational drug design process. For this purpose, we introduce the docking algorithm PLANTS (Protein–Ligand ANT System), which is based on ant colony optimization, one of the most successful swarm intelligence techniques. We study the effectiveness of PLANTS for several parameter settings and present a direct comparison of PLANTS’s performance to a state-of-the-art program called GOLD, which is based on a genetic algorithm and frequently used in the pharmaceutical industry for this task. Last but not least, we also show that PLANTS can make effective use of protein flexibility giving example results on *cross-docking* and *virtual screening* experiments for *protein kinase A*.

Keywords Ant colony optimization · Protein–ligand docking · Cross-docking · Virtual screening

1 Introduction to the protein–ligand docking problem

The pharmacological effect of a drug relies in many cases on the formation of a complex between a small chemical compound, the drug called *ligand* in the following, and a macro-

This article is based on a paper that won the best paper award at ANTS 2006, the 5th International Workshop on Ant Colony Optimization and Swarm Intelligence held in Brussels, Belgium, 2006. This article includes new types of experiments and also the possibility of considering flexibility of protein side-chains.

O. Korb · T.E. Exner
Theoretische Chemische Dynamik, Universität Konstanz, Konstanz, Germany

O. Korb
e-mail: Oliver.Korb@uni-konstanz.de

T.E. Exner
e-mail: Thomas.Exner@uni-konstanz.de

T. Stützle (✉)
IRIDIA, CoDE, Université Libre de Bruxelles, Brussels, Belgium
e-mail: stuetzle@ulb.ac.be

molecule, the protein, whose malfunction is provoking a disease. For the rational design of new drugs against a target protein, the specific interactions between both complex partners must be correctly predicted and optimized. This task is notoriously time-consuming and expensive taking up to 15 years (Müller 2003) and costing several hundred million euros.

This drug design process pursued by major pharmaceutical companies begins with the identification of a suitable protein target, where a potential drug could interfere to fight a disease. For this target, specific assays are developed, which are then used in *high-throughput screening* experiments to test the biological activity of large databases of possible drug candidates. Molecules with high affinity, so-called lead structures, are then chemically varied (lead optimization cycle) and the most potent ones of the resulting candidates are transferred to the preclinical and finally the clinical test phase.

Today, most of these steps are guided by computer experiments. Very large databases, ranging in size from thousands to millions of ligands, can be tested by ligand- or structure-based *virtual screening* techniques and different lead-optimization routes can be tried without ever synthesizing most of the compounds. Structure-based approaches, like the one presented in this publication, start with a known 3D structure of a protein. These structures are obtained by experimental techniques like X-ray crystallography or NMR-spectroscopy and are, for example, publicly available from the *protein data bank* (PDB) (Berman et al. 2000). Then, a docking algorithm tries to solve the *pose prediction* problem. This problem consists in finding the correct orientation and conformation (that is, the atoms' spatial arrangement resulting from rotations about single-bonds) of the ligand within the *a priori* known active site of the protein. In the context of *virtual screening*, additionally different ligands must be ranked correctly according to their binding affinity.

The underlying so-called protein–ligand docking problem (PLDP) was first formulated by Fischer using the famous lock-and-key metaphor (Fischer 1894): The key (ligand) must fit exactly into the lock (protein) to open the door (pharmacological effect). But because molecules are not rigid objects, this description is too limited and at least the flexibility of the ligand must be taken into account, which is done by almost all state-of-the-art docking algorithms. But even the protein conformation can adapt to an incoming ligand, leading to the so-called *induced fit* model (Koshland 1973). In particular, if the protein structure that has been determined without any bound ligand, the so-called *apo form*, is used, large rearrangements at the active site of the protein may be observed. Therefore, whenever possible, a 3D-structure of a protein complexed with a known ligand, a *holo form*, is used under the assumption that a potential new drug would bind to an extremely similar protein conformation. The validity of this assumption is apparent in many cases when comparing different *holo forms* of the same protein cocrystallized with different ligands whose structures are stored in the PDB.

2 Computational approaches to the docking problem

A large variety of approaches for solving the PLDP has been proposed. Broadly, these can be classified as *fragment-based* methods, which split the ligand into different parts and perform an incremental construction, as *stochastic optimization* methods, or as *multiconformer docking* approaches (see Taylor et al. 2002 and references therein). In this paper, we focus on the second class. Stochastic optimization methods are based on the formulation of the PLDP as an optimization problem and on the use of stochastic search methods for searching optimal solutions. The optimization problem is defined by an objective function f and the search space. In the context of the PLDP, the objective function is usually called *scoring*

function and it estimates the binding energy between the protein and the ligand; the search space is given by the degrees of freedom of the protein and the ligand. Given these two components, the PLDP can be seen as the problem of searching for the values of the degrees of freedom that globally minimize the scoring function.

In most approaches, the protein structure is kept rigid. In this case, the task is to find the best possible orientation of the ligand with respect to the protein by changing the ligand's translation and rotation as well as changing the torsion angles of single bonds in the ligand that are not part of a ring system. In the optimization version of the PLDP, these degrees of freedom correspond to continuous variables that define the transformation applied.

Hence, the task is to find optimal values for the ligand's 3 translational, 3 rotational and r_l torsional degrees of freedom, describing the rotations of single bonds. Thus, the total number of variables, that is, the dimension of the optimization problem, equals $n = 6 + r_l$. In more advanced methods, like the approach presented here, small changes in the conformation of the protein can be modeled by introducing flexibility to specific amino-acid side-chains in the active site. This adds one torsional degree of freedom for each rotatable bond in the flexible side-chains. Therefore, the problem dimension becomes $n = 6 + r_l + r_p$ for a ligand with r_l and a protein with r_p torsional degrees of freedom.

A wide spectrum of optimization strategies has been proposed for finding the global minimum energy structure, which is expected to correspond to the experimentally determined complex structure. These include genetic algorithms that are used in the programs GOLD (Jones et al. 1997) and AutoDock (Morris et al. 1998), Monte Carlo minimization in the programs ICM (Abagyan et al. 1994) and QXP (McMartin and Bohacek 1997), and simulated annealing, evolutionary programming, and tabu search in PRO_LEADS (Baxter et al. 1997).

Recent studies (Kellenberger et al. 2004; Kontoyianni et al. 2004) compared different docking tools on a large test set of experimentally determined complex structures. They reported success rates of 30 to 60%, where the success rate is defined as the percentage of complexes for which the predicted structure with the lowest energy is very close (root mean square deviation within 2.0 Å) to the experimentally determined structure. Therefore, it can be concluded that a universal docking tool that has excellent predictive capabilities across many complexes is not available at the moment. When we consider, as we do in this paper, stochastic optimization methods, this lack of performance can be attributed to the *scoring problem* and the *sampling problem*. Currently, there exists no scoring function capable of performing correct measurements for all given input structures. But even if there was a perfect scoring function, there would still be the problem that there is no guarantee that the correct binding mode of the ligand, that is, the ligand conformation as observed in the crystal structure, is actually found by the sampling algorithm. On the one hand, this can be caused by the fact that the native complex structure is not accessible with the limited flexibility of the protein accounted for in the algorithms. On the other hand, introducing additional degrees of freedom may result in additional scoring function failures or make the problem more difficult due to the increased dimension of the search space.

3 PLANTS

We have developed Protein–Ligand ANT System (PLANTS), the first ACO algorithm for tackling the PLDP. PLANTS is a stochastic search algorithm that treats the ligand and the protein as flexible, which means that there are $6 + r_l$ degrees of freedom for the ligand and r_p torsional degrees of freedom for the protein as described above. Even if no flexible side-chains are specified, PLANTS partially considers the flexibility of the protein by allowing




Fig. 1 This figure illustrates the degrees of freedom for the docking problem. It shows the origin of the ligand's coordinate system as a sphere. The *large arrows* indicate the ligand's translational degrees of freedom; they also are the axes of rotation. The *small arrows* mark the ligand's rotatable bonds and also the rotatable bonds of a single protein side-chain (*upper right corner*), which emanates from the protein surface that is shown in the background

for the optimization of the positions of hydrogen atoms that could be involved in hydrogen bonding; this results in $r_p = r_{\text{don}}$ torsional degrees of freedom, where r_{don} is the number of rotatable hydrogen bond donor groups (OH- and NH_3^+ -groups) available in the protein binding site. Both the ligand's and the protein's degrees of freedom are illustrated in Fig. 1.

The resulting minimization problem to be solved by PLANTS can be formulated as

$$\min_{\vec{x} \in \mathbb{R}^n} f(\vec{x}) : \mathbb{R}^n \rightarrow \mathbb{R}. \quad (1)$$

In this equation, f is the objective function, also called the *scoring function*, $\vec{x} = [x_1, \dots, x_n]^T \in \mathbb{R}^n$ represents the ligand's and the protein's degrees of freedom and $n = 6 + r_l + r_p$ is the problem dimension. Hence, the problem is a continuous optimization problem, where the angles can be in the domain $[0^\circ, 360^\circ)$ and the intervals of the domains for the translational degrees are defined by the size of the protein's binding site.

Pheromone model The space searched by PLANTS is defined by the ligand's translational, rotational and torsional degrees of freedom as well as the protein's torsional degrees of freedom. Since ACO is initially designed for solving combinatorial problems, we decided to discretize the continuous variables and to apply existing ACO algorithms, in particular *MAS-MIN* Ant System (*MMAS*) (Stützle and Hoos 2000), to the problem. The discretization uses for each of the three translational degrees of freedom an interval length of 0.1 \AA , while for the three rotational degrees of freedom and all torsional degrees of freedom an interval length of 1° was taken, resulting in 360 values for the latter. The number of values for each translational degree of freedom depends on the diameter of the pre-defined size of the binding site. For the complexes studied here, this discretization results in 120 to 400 discrete points for each dimension x , y and z . Each degree of freedom i has associated a pheromone vector $\vec{\tau}_i$ with as many entries as values resulting from the discretization. A pheromone trail τ_{ij} then refers to the desirability of assigning the value j to degree of freedom i .

ACO algorithm The (artificial) ants in PLANTS construct solutions by selecting one value for each degree of freedom taking into account the pheromone values. In an earlier paper (Korb et al. 2006), we have shown that heuristic information was not essential for high

performance and, hence, we removed it in this more recent version of PLANTS. In the solution construction, the order of the degrees of freedom is arbitrary, since each degree of freedom is treated independently of the others. For each degree of freedom i , the probability p_{ij} of assigning the value j to an ant is computed as

$$p_{ij} = \frac{\tau_{ij}}{\sum_{l=1}^{n_i} \tau_{il}}. \quad (2)$$

In this equation, n_i is the number of values for degree of freedom i . As usual in \mathcal{MMAS} , after each iteration only one solution deposits pheromone; in PLANTS, this is the best solution generated in the current iteration, s^{ib} . The pheromone update is defined as

$$\tau_{ij}(t+1) = (1 - \rho)\tau_{ij}(t) + I_{ij}^{ib}(t)\Delta\tau^{ib}(t), \quad (3)$$

where

$$\Delta\tau^{ib}(t) = \begin{cases} |f(s^{ib})|, & \text{if } f(s^{ib}) < 0, \\ 0, & \text{otherwise,} \end{cases} \quad (4)$$

$f(s^{ib})$ is the value of the scoring function for s^{ib} , and ρ is a parameter steering the pheromone evaporation. For a translational degree of freedom i , $I_{ij}^{ib}(t)$ is one, if s^{ib} assigned to i a value in $\{j-1, j, j+1\}$; for rotational and torsional degrees of freedom, $I_{ij}^{ib}(t)$ is one if a value in $\{j-2, j-1, j, j+1, j+2\} \bmod n_i$ was taken; otherwise it is zero. In this way, for each degree of freedom not only one single value receives some pheromone but also the neighboring values. This can be useful to increase the search diversification. The choice of (4) is based on the following reasoning: strongly negative values of our scoring function indicate high affinity, that is, the larger the absolute value the better; positive values would actually correspond to negative affinity and, hence, they do not receive any positive feedback. In the latter case, no pheromone is deposited in an iteration. In PLANTS, \mathcal{MMAS} 's lower and upper pheromone trail limits, τ_{\min} and τ_{\max} , are set as follows: $\tau_{\max} = |f(s^{db})|/\rho$, where $f(s^{db})$ is the best scoring function value found since the last search diversification (the search diversification used is explained below); the values for τ_{\min} are calculated for each degree of freedom i as

$$\tau_{\min_i} = \frac{\tau_{\max} \cdot (1 - \sqrt[n]{p_{\text{best}}})}{(n_i - 1) \cdot \sqrt[n]{p_{\text{best}}}}, \quad (5)$$

following the formulas given in Stützle and Hoos (2000). In this equation, n is the problem dimension, that is, the number of degrees of freedom, n_i is the number of values that resulted from the discretization of degree of freedom i , and p_{best} is the probability of the best solution to be reconstructed assuming that the colony has already converged (Stützle and Hoos 2000).

Local search As in most applications of ACO algorithms to \mathcal{NP} -hard problems, we also improve candidate solutions by a local search algorithm (Dorigo and Stützle 2004). However, differently from most of such combinations that have been proposed earlier, in PLANTS the ants and the local search algorithm search in two different search spaces: while the ants work on the discretized search space, the local search works in the continuous search space defined by the translational, rotational and torsional degrees of freedom of the ligand and the protein. In fact, the resulting, very high performing algorithm for the PLDP indicates that such a combination may also be interesting for other nonlinear continuous optimization problems.

We use the simplex local search algorithm described by Nelder and Mead (NMS) for nonlinear, continuous function optimization (Nelder and Mead 1965). A simplex is a polytope of $n + 1$ vertices in an n -dimensional space. The NMS algorithm transforms the points of a given starting simplex by using the operations reflection, expansion and contraction until the fractional range from the highest to the lowest point in the simplex with respect to the function value is less than a tolerance value, which we choose as 0.01, or a maximum number of function evaluations is reached (we set this limit to 5000 function evaluations); for details on the algorithm we refer to (Press et al. 1992). As the parameter settings in the NMS algorithm we use $\Delta_{\text{trans}} = 2\text{\AA}$, $\Delta_{\text{rot}} = 90^\circ$ and $\Delta_{\text{tors}} = 90^\circ$ for the construction of the initial simplex.

In PLANTS, all ants improve their solution by applying NMS. The starting simplex is defined by the initial, unmodified solution of each ant and the other n by adding the specific offsets to each variable. Note that the NMS algorithm does not necessarily stop in a local optimum and, hence, the final solution could be further improved by re-applying NMS. In fact, we do so and use NMS again to further improve the solution of the iteration-best ant. This refinement local search, which uses the same parameter settings as given above, is restarted until the improvement in the scoring function obtained by one application of NMS is less than 0.2.

Algorithmic outline An algorithmic outline of PLANTS is given in Algorithm 1. Essentially, PLANTS follows the typical scheme of solution construction, local search, and pheromone update that is followed in virtually all ACO algorithms for \mathcal{NP} -hard combinatorial optimization problems (Dorigo and Stützle 2004). Some additional details are explained next.

The procedure `InitializeParametersAndPheromones` sets the initial parameter values, sets up the ligand and the protein for processing, and initializes the pheromone trails to a very large value such that after the first iteration all pheromone trails correspond to τ_{max} .

Algorithm 1 PLANTS

```

InitializeParametersAndPheromones()
for  $i = 1$  to  $\text{max\_iterations}$  do
    for  $j = 1$  to  $\text{ants}$  do
         $s_j \leftarrow \text{ConstructSolution}()$ 
         $s_j^* \leftarrow \text{LocalSearch}(s_j)$ 
         $M \leftarrow M \cup s_j^*$ 
    end for
     $s^{ib} \leftarrow \text{GetBestSolution}()$ 
     $s^{ib} \leftarrow \text{RefinementLocalSearch}(s^{ib}, 0.2)$ 
     $M \leftarrow M \cup s^{ib}$ 
     $\text{UpdatePheromones}(s^{ib})$ 
    if  $\text{diversificationCriteriaMet}$  then
         $\text{ApplySearchDiversification}()$ 
    end if
end for
return  $M$ 
  
```

The number of iterations for PLANTS is determined by the formula

$$\max_iterations = \sigma \cdot \frac{10}{m} \cdot (100 + 50 \cdot lrb + 5 \cdot lha), \quad (6)$$

where σ is a parameter used for scaling the number of iterations, m is the ant colony size, lrb is the number of rotatable bonds, and lha the number of non-hydrogen atoms in the ligand. Hence, the maximum number of iterations depends on the ligand: very flexible and large ligands get more search time than rigid and small ones.

The procedures **ConstructSolution** and **LocalSearch** implement the solution construction and the NMS local search algorithm, respectively. Once for each ant a solution is generated, the iteration-best solution s^{ib} is determined by the procedure **GetBestSolution**. The procedure **RefinementLocalSearch** then applies the additional local searches to s^{ib} , as described in the previous subsection. Next, the procedure **UpdatePheromones** manages the pheromone update and checks the pheromone trail limits.

PLANTS also applies additional search diversification in procedure **ApplySearchDiversification**. Search diversification consists, in the first place, of pheromone trail smoothing as proposed in Stützle and Hoos (2000), using a smoothing factor of 0.5. If three smoothings have been done, all pheromone trails are erased and set to their initial value. In fact, this latter diversification corresponds to a complete restart of the algorithm. Search diversification is invoked each time more than 10 consecutive iteration-best solutions differ by less than $0.02 \cdot |f(s^{db})|$, where s^{db} is the best solution found since the last search diversification.

Once PLANTS terminates, the set of all solutions (M) returned by the procedures **LocalSearch** and **RefinementLocalSearch** is output; this set M also contains the overall best solution found by PLANTS. The solutions in set M then undergo a post-processing phase.

Post-processing Once PLANTS finishes its search, all solutions that were generated are post-processed to output a number of diverse, high-quality conformations. This is done by first sorting all the solutions in M according to increasing scoring function values. Then, PLANTS extracts a specified number of ligand structures, typically 10, such that the minimal root mean square deviation (RMSD) between any of the extracted structures is larger than 2 Å. These solutions can then be used for rescoring with other, more advanced scoring functions, which may be computationally too demanding to be directly used in the optimization process. Experimentally, it has been found that rescoring can increase the chance of finding a ligand conformation that is similar to the experimental binding mode. Examples of rescoring protocols can, for example, be found in Verdonk et al. (2003).


Empirical scoring functions In PLANTS, two empirical scoring functions are available: a modified *piecewise linear potential* (PLP) version (Gehlhaar et al. 1995; Verkhivker 2004) and a new scoring function, called CHEMPLP, that is assembled by combining parts of already published ones (Verdonk et al. 2003; Gehlhaar et al. 1995).

The PLP scoring function, f_{PLP} , employed in PLANTS is based on the ones described in Gehlhaar et al. (1995), Verkhivker (2004) using distance-based potentials only and it has the form

$$f_{PLP} = f_{plp} + f_{tors-lig} + f_{clash-lig} + 0.3 \cdot f_{score-prot} - 20.0. \quad (7)$$

The first part, f_{plp} , mainly models steric interactions between the protein and the ligand, considering metal ions in the protein binding site and the occlusion of polar atoms by non-polar ones by distance-based potentials. The general functional form of this term can be seen in Fig. 2a, where parameter r is the Euclidean distance between a ligand and a protein

Fig. 2 Illustration of some empirical scoring function terms. **a** Intermolecular *piecewise linear potential* (PLP). **b** Hydrogen bond geometry descriptors. The parameters are explained in the text



atom and $PLP(r)$ is the resulting potential value. Parameters A to F define the form of the potential and vary depending on the protein and ligand atom type (Gehlhaar et al. 1995; Verkhivker 2004). The intramolecular ligand scoring function consists of a simple clash term ($f_{\text{clash-lig}}$), which avoids the ligand's atoms to come too near, and a torsional potential ($f_{\text{tors-lig}}$) as described in Clark et al. (1989). For intramolecular protein-interactions ($f_{\text{score-prot}}$), the same distance-based potentials as in PLP together with an intra-side-chain clash term is used.

Scoring function CHEMPLP, denoted as f_{CHEMPLP} has the following functional form:

$$f_{\text{CHEMPLP}} = f_{\text{plp}} + f_{\text{chem-hb}} + f_{\text{tors-lig}} + f_{\text{clash-lig}} + 0.3 \cdot f_{\text{score-prot}} - 20.0. \quad (8)$$

The first part (f_{plp}) of the intermolecular score uses the above described version of the PLP scoring function, although with different parameter settings. The second part, ($f_{\text{chem-hb}}$) considers the hydrogen bonding and metal-acceptor interactions between the protein and the ligand as done in GOLD's CHEMSCORE implementation (Verdonk et al. 2003). Figure 2b shows an illustration of the parameters used to define the hydrogen bond geometry. The Euclidean distance r between the hydrogen bond donor atom D (donor) and a hydrogen bond acceptor atom A (acceptor) as well as the donor angle α and the acceptor angle β influence the strength of the hydrogen bond depending on their deviation from a given set of ideal values. This part has been modified to differentiate between charged and neutral hydrogen bonds and to allow for the scoring of weak CH–O interactions (Verdonk et al. 2004). The intramolecular ligand and protein terms are the same as described in the PLP case.

Finally, a penalty term is added to both scoring functions if the ligand's reference point falls outside the predefined binding site of the protein. The development of these empirical scoring functions was, together with the design of the search algorithm, a major contribution of our work; the development of the scoring function will be covered in a forthcoming publication.

4 Parameter optimization

In an earlier paper on PLANTS (Korb et al. 2006), the influence of the number of ants, the pheromone evaporation ρ , the heuristic information and the factor σ , which scales the number of iterations, was investigated. It turned out that using 20 ants in conjunction with a setting of $\rho = 0.25$ seems to result in overall good performance across different settings of σ . Heuristic information was found to have no significant impact on performance and, hence, we removed it from further consideration in our current version of PLANTS. These

earlier tests were all carried out using a setting of 0.9 for parameter p_{best} , which directly affects the lower pheromone trail limit τ_{min} and, hence, also the explorative power of the algorithm. In this paper, we study the parameters that influence the diversification behavior of the algorithm. In this context, the parameters p_{best} and the evaporation rate ρ are of special interest.

Two subsets comprising 15 complexes (for scoring function CHEMPLP) and 11 complexes (for scoring function PLP) from the CCDC/Astex dataset (Nissink et al. 2002) with ligands having between two and ten rotatable bonds have been used for the parameter variation. These subsets were removed from the final test set used for the large-scale tests described in the next section. We varied the parameters σ , ρ and p_{best} considering three to seven values for each (the values tested are indicated in Fig. 3), which resulted in 84 distinct parameter configurations. On each complex, PLANTS was run for 25 independent trials. We measured for each configuration the average success rate, the average computation time and the average number of function evaluations. The success rate is defined as the percentage of complexes for which the top-ranked docking solution is within an RMSD of 2.0 Å of the experimentally determined binding mode as given in the CCDC/Astex dataset (Nissink et al. 2002). The computation times in this section are given in seconds on a single core of a Dual Core AMD Opteron 870, 2.0 GHz CPU; protein setup time (6 s on average) and ligand setup time (0.01 s on average) are excluded.

Figure 3 gives a graphical analysis of the results obtained for the variation of parameter ρ (Fig. 3a) and parameter p_{best} (Fig. 3b) when scoring function CHEMPLP is used for the docking experiments. The same parameter optimization was also carried out for scoring function PLP; since the overall trends for PLP are the same as for CHEMPLP, we only discuss in what follows the results obtained for scoring function CHEMPLP.

The data points of Fig. 3a and Fig. 3b obviously fall into three distinct clusters with respect to the docking time in dependence of the value of σ (the values 0.25, 0.5, and 1.0 have been tested), which scales the number of iterations carried out by the algorithm. The three clusters are separated by dashed lines in the plots (left cluster: $\sigma = 0.25$; middle cluster: $\sigma = 0.5$; right cluster: $\sigma = 1.0$).

The parameter ρ , as was also observed in Korb et al. (2006), has a rather direct impact on the computation time and, in part, also on the success rate. This can be seen by the fact that the points in Fig. 3a fall, for each value of σ , into roughly four smaller clusters, identified by a same value of ρ . For $\sigma = 0.25$, the setting $\rho = 0.1$ results in somewhat higher success rates, while for the other values of σ the differences are less pronounced.

Concerning the parameter p_{best} , no clear trend is apparent and the conclusion is that the performance of PLANTS is not very sensitive with respect to p_{best} .

For the sake of having a same set of parameter values across different values of σ , we use a setting of $\rho = 0.15$ and $p_{\text{best}} = 0.5$ in the case of scoring function CHEMPLP as well as $\rho = 0.15$ and $p_{\text{best}} = 0.1$ for scoring function PLP when performing tests on a large set of benchmark instances. Hence, all the experiments presented in the next section differ only in the setting of parameter σ , which allows to directly trade between computation time and success rates.

5 Pose prediction

The *clean list* of the comprehensive CCDC/Astex dataset (Nissink et al. 2002) has been used for the validation of PLANTS. From these 224 complexes, 11 include covalently bound ligands and were removed, because PLANTS does not provide a covalent docking functionality at the moment. Additionally, also the 15 complexes used during the parameter optimization process described in Sect. 4 were removed to avoid a bias in favor of PLANTS. Hence, our

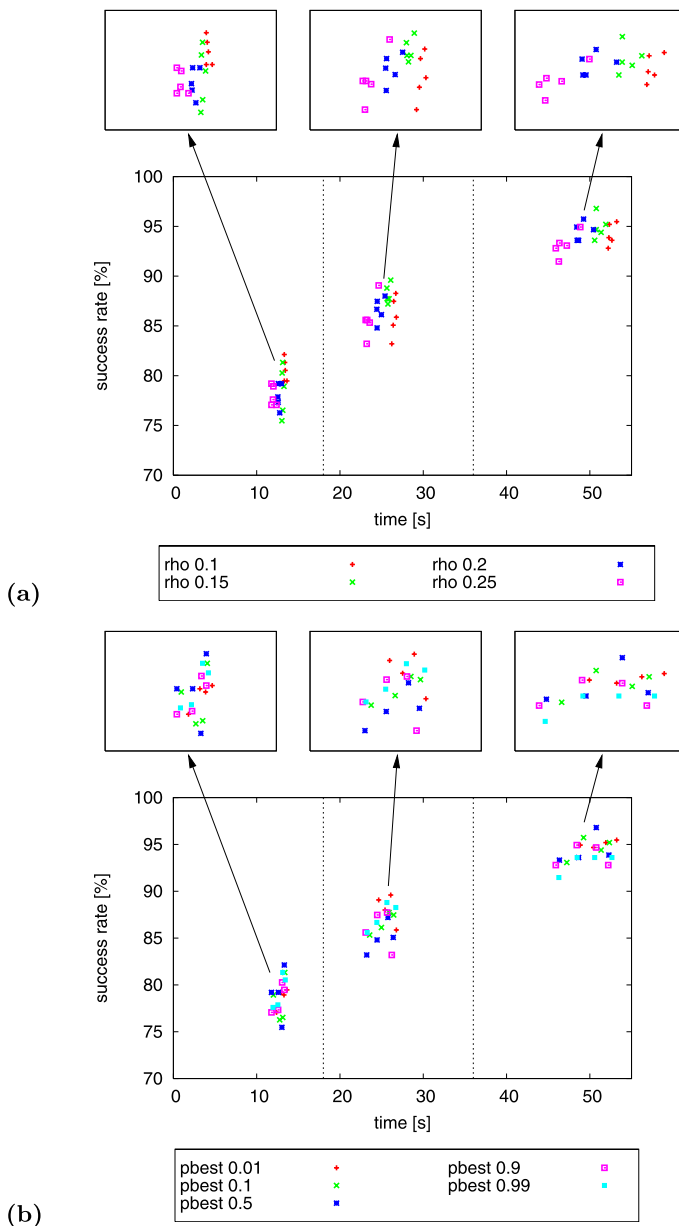


Fig. 3 Influence of different settings for **a** evaporation factor ρ and **b** parameter p_{best} with respect to the average success rate and docking time for scoring function CHEMPLP. For the sake of clarity, the data points for $p_{\text{best}} = 0.3$ and $p_{\text{best}} = 0.7$ are not shown as they were anyway not part of the optimal parameter set. For further explanations see the text. Note that the zooms of the individual clusters differ in their x/y -ratio slightly from the original plot to improve the visibility of the distribution of the points

Table 1 Results on 198 complexes of *clean list^{nc}* for PLANTS (scoring function CHEMPLP and PLP) and GOLD (scoring function GOLDScore) for selected parameter settings averaged over 25 independent experiments. For all tests a colony size of 20 ants was used. Standard deviations for the success rates are given in parentheses

PLANTS _{CHEMPLP}					
σ	success rate (%) up to rank			Time (s)	Eval. (10^6)
	1	3	10		
0.25	62.93 (2.58)	74.71 (1.95)	80.75 (2.01)	27.85	1.02
0.40	65.78 (2.09)	78.14 (1.48)	84.30 (1.34)	43.98	1.61
1.00	71.94 (1.84)	83.94 (1.48)	89.62 (1.46)	110.07	3.99
2.00	73.80 (1.37)	86.40 (1.35)	92.04 (1.43)	217.04	7.92
PLANTS _{PLP}					
σ	success rate (%) up to rank			Time (s)	Eval. (10^6)
	1	3	10		
0.50	62.44 (2.00)	75.52 (1.64)	83.19 (1.65)	22.25	1.26
1.00	64.32 (1.51)	78.55 (1.13)	86.16 (1.46)	45.24	2.50
2.00	65.25 (0.93)	79.76 (1.33)	87.72 (1.44)	89.64	4.98
5.00	66.08 (1.08)	80.57 (1.16)	89.43 (1.35)	218.97	12.41
GOLD _{GOLDScore}					
autoscale	success rate (%) up to rank			Time (s)	Eval. (10^6)
	1	3	10		
0.10	65.94 (1.56)	72.61 (1.26)	77.19 (1.51)	43.47	n.a.
0.30	67.94 (1.90)	74.16 (1.88)	80.08 (2.12)	118.84	n.a.
1.00	72.12 (1.57)	76.81 (1.53)	81.31 (1.23)	318.02	n.a.

test set consists of 198 non-covalently bound complexes that we call *clean list^{nc}*. The number of rotatable bonds of the ligands in *clean list^{nc}* ranges from 0 to 28. For all experiments, the spherical binding site as defined in the CCDC/Astex dataset was used to determine the search space for the ligand's translational degrees of freedom. Before each docking run, the ligand structures were randomized with respect to the translational, rotational, and torsional degrees of freedom, that is, the ligands' conformations as observed in the crystal structures were changed to random conformations to avoid biasing the results. The best parameter settings determined in the last section were then tested on the whole *clean list^{nc}*. The applied parameter setting as well as the success rate for the (i) top-ranked solution, (ii) solutions up to rank 3, and (iii) solutions up to rank 10 (ranks with respect to the solutions in the order as returned by the post-processing—a success is obtained if among these highest ranked ligands one has an RMSD lower than 2 Å compared to the experimentally determined complex structure) and the average docking time along with the number of scoring function evaluations are presented in Table 1 (see upper and middle part marked with PLANTS). All calculations were carried out on a Pentium 4 Xeon, 2.8 GHz CPU.

For PLANTS_{CHEMPLP}, the success rates for the top-ranked solutions range from about 62.9% at average docking times of approximately 28 s ($\sigma = 0.25$) to 73.8% at average docking times of 217 s ($\sigma = 2$) for each complex. Considering solutions up to rank 3 and 10 even higher success rates of up to 92% can be achieved. This is especially interesting for *vir-*

tual screening applications where rescoring of these structures with another scoring function could in principle identify the correct pose. In dependence of the available computational resources, the time frame and the size of the ligand database to be screened (which can be in the order of thousands to millions of compounds), a potential user may choose one of the parameter configurations presented here. As a guideline we suggest to use a setting of 20 ants, $\sigma = 1$, $\rho = 0.15$, and $p_{\text{best}} = 0.5$ as the standard setting for PLANTS.

Using the other scoring function available in PLANTS, PLP, the algorithm reaches success rates ranging from about 62% at average search times of around 22 s to about 66% at average docking times of 219 s per complex considering the first-ranked solutions only. For low computation time limits, the results of PLANTS_{PLP} are reasonably close to those obtained by using PLANTS_{CHEMPLP}. This suggests that PLANTS profits from the additional hydrogen bonding terms available in CHEMPLP especially for higher computation times. In fact, PLANTS_{PLP} seems to reach limiting behavior at success rates of about 66%, given that an increase of σ from two to five has only a negligible effect on increasing the success rate. This trend can be confirmed by statistical tests on the significance of the differences in the success rates. Using the binomial test with significance level $\alpha = 0.05$ for PLANTS_{CHEMPLP} and PLANTS_{PLP} for $\sigma = 0.4$ and $\sigma = 1.0$, respectively (this results in similar average computation times), no significant difference between the two could be detected; however, for a setting of $\sigma = 2.0$ and $\sigma = 5.0$ for PLANTS_{CHEMPLP} and PLANTS_{PLP}, respectively, the null hypothesis had to be rejected at the $\alpha = 0.05$ level. A similar trend is true when looking at the results up to rank three and ten. In this case, for short computation times PLANTS_{PLP} sometimes even gives slightly higher success rates for comparable computation times, while for the longer computation times PLANTS_{CHEMPLP} becomes preferable. (When comparing the same settings of σ for PLANTS_{CHEMPLP} and PLANTS_{PLP}, only the difference between the success rates up to rank three for the longer computation times are statistically significantly different.)

For comparison, the lower part of Table 1 shows the results on the same test set for GOLD (Genetic Optimization for Ligand Docking) (Jones et al. 1997; Verdonk et al. 2003), a state-of-the-art docking program that is frequently used in the pharmaceutical industry. The genetic algorithm optimizes mappings of possibly corresponding hydrogen bonding atoms and hydrophobic groups in the protein and the ligand as well as the ligand's torsional degrees of freedoms and the orientation of hydrogen bond donor groups in the protein. In each GA-step, a matching procedure tries to minimize the distance between the mapped fitting points. Additionally, problem-specific knowledge like torsion-angle libraries for the ligands' rotatable bonds as well as a cavity detection algorithm for restricting the search space for possible ligand placements is employed.

Ideally, when comparing the two approaches, both programs should use the same scoring function to focus the comparison on the impact of the search algorithms (or, if the scoring functions are of interest, these are tackled using the same search algorithm). However, integrating either of the scoring functions into the other program package is first of all a non-trivial, time-consuming task and second it inherently has the problem that the reimplementation and the original may not correlate perfectly. Additionally, changing the scoring function requires the search algorithm to be adapted and fine-tuned as well, making it possibly very different from the original implementation. Hence, we rather opted for comparing the two programs' performance as a whole. There are good reasons for doing so. Firstly, this allows to compare the performance of PLANTS to a recognized state-of-the-art package for the same task; secondly, this comparison is also relevant for practice, where the algorithms would be used out-of-the-box for a particular docking task.

For the experiments presented in this section, GOLD version 3.0.1 has been employed. We tested three settings for parameter *autoscale*, which automatically chooses appropriate

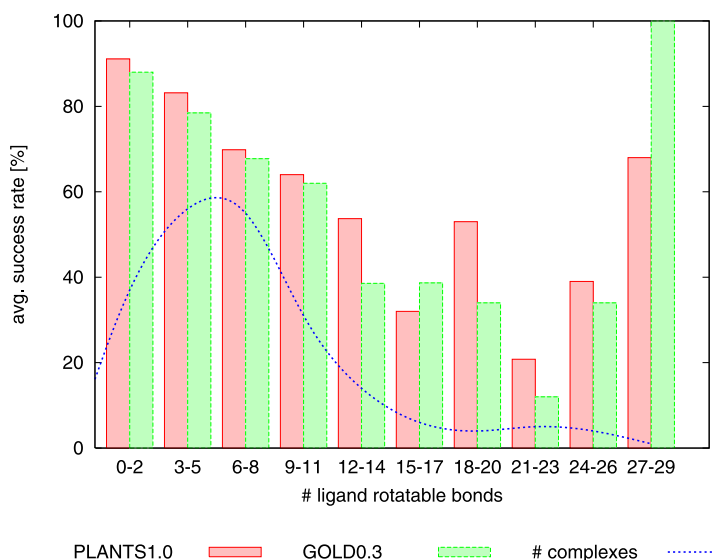


Fig. 4 Average success rates in dependence of the number of ligand rotatable bonds for PLANTS_{CHEMPLP} ($\sigma = 1$) and GOLD_{GOLDScore} (*autoscale* = 0.3). Additionally, the number of complexes for each interval of ligand rotatable bonds is shown by the dotted line

GA search settings in dependence of the given protein binding site and the ligand molecule. The maximum number of GA runs per ligand was set to 10, *early termination* and *cavity detection* were activated and *Goldscore* (GOLD_{GOLDScore}) was used as the scoring function.

When comparing the success rates obtained by PLANTS_{CHEMPLP} to those of GOLD_{GOLDScore} at similar or slightly lower average computation times given in Table 1 (that is, settings of $\sigma \in \{0.4, 1.0, 2.0\}$ for PLANTS_{CHEMPLP}), PLANTS_{CHEMPLP} reaches always higher success rates with the only exception being the very minor difference in favor of GOLD_{GOLDScore} for the top-ranked solution at a setting of $\sigma = 0.4$ for PLANTS_{CHEMPLP}. (The binomial test at $\alpha = 0.05$ for equality of success rates was rejected in favor of PLANTS_{CHEMPLP} for all pair-wise comparisons of the success rates up to rank three and ten.) Hence, this clearly shows that PLANTS's performance is at least comparable to that of GOLD, which is a very encouraging result. Finally, more detailed information on the success rate of PLANTS_{CHEMPLP} versus GOLD_{GOLDScore} concerning dependence of the number of rotatable ligand bonds is given in Fig. 4, which again confirms the competitiveness of PLANTS.

6 Protein flexibility

The pose prediction tests carried out in the previous section for the *clean list*^{nc} were ideal test cases in the sense that the ligands were docked back into their native protein structures. When performing a *virtual screening* experiment, however, usually only a single protein structure is available where all biologically active ligands should fit into with the correct conformation. As already mentioned in the introduction, upon ligand binding the protein may undergo small or sometimes even large rearrangements, in which case a single rigid protein structure may not be able to reproduce the poses of all ligands correctly. In this

section, we deal with exactly this type of problem and examine the influence of protein flexibility during docking and *virtual screening*. In fact, the importance of protein flexibility is now widely recognized and very recent versions of some docking algorithms like GOLD include the treatment of *induced fit* effects.

6.1 Cross-docking

As a first step, we test PLANTS with respect to its ability to reproduce binding modes of different ligands also in non-native protein structures of the same target, a type of experiment called *cross-docking*. The *cross-docking* experiment was carried out for *protein kinase A* (PKA), which is a promising target in breast-cancer research. For this experiment, we have used eight PKA protein structures (seven *holo forms* and one *apo form*). We used PDB-entries 1Q8T (ligand *y27*), 1Q8U (ligand *h52*), 1Q8W (ligand *m77*), 1STC (ligand staurosporine, *stu*), 1YDR (ligand *iqp*), 1YDS (ligand *iqs*), 1YDT (ligand *iqb*) (these seven are *holo forms*, that is, proteins with a complexed ligand), and 1JLU (*apo form* of the protein, no ligand). The superimposition of all the protein backbones is illustrated in Fig. 5; each protein backbone is indicated as one tube. The bold circle in Fig. 5 indicates a clash between one ligand (staurosporine from PDB-entry 1STC) with protein residue PHE327 for all non-native protein structures—this residue is shown for all eight protein structures in the lower left corner. Therefore, to allow for a correct pose prediction of the ligand staurosporine also in non-native protein structures, at least the side-chain of residue PHE327 needs to be treated flexible.

Starting from these observations, two *cross-docking* experiments have been carried out. In the first case, the seven ligands have been docked into all eight protein structures considering no protein flexibility, while in the second case the side-chain of residue PHE327 was treated as flexible. In addition to the ten degrees of freedom resulting from rotatable donor groups in the protein, which are also considered in the rigid protein case, two more degrees of freedom are introduced by the flexible side-chain. However, the conformational space induced by these two degrees of freedom has a greater impact on the docking solutions than the one induced by the ten rotatable donor groups. Due to structural deviations in the superimposition of the protein backbones, a docking success criterion of RMSD <2.5 Å for

Fig. 5 Superimposition of the eight PKA protein structures (the protein backbones are illustrated as tubes). In the *lower left corner* the side-chain conformations of residue PHE327 for each protein structure are shown (for the sake of clarity all other residues are hidden). Ligand staurosporine (shown in the middle along with all other ligands) clashes with this residue in all non-native structures (marked with a **bold circle**)




Table 2 PLANTS *cross-docking* results. The table shows the success rate in percent obtained for docking the seven ligands into rigid (upper part) and partially flexible (lower part) protein structures. The bold entries on the diagonal correspond to the docking of a ligand into its native protein structure. All results are averaged over 25 independent runs

docking into rigid protein structures								
ligand	1Q8T	1Q8U	1Q8W	1STC	1YDR	1YDS	1YDT	1JLU
y27	100	80	100	0	100	100	100	100
h52	100	100	100	60	88	100	100	8
m77	100	96	100	20	100	100	100	88
stu	0	0	0	100	0	0	0	0
iqp	4	0	100	96	100	100	92	48
iqs	44	0	96	12	100	100	100	60
iqb	48	76	32	0	32	60	72	20
avg.	56.57	50.29	75.43	41.14	74.29	80.00	80.57	46.29
docking into flexible protein structures								
ligand	1Q8T	1Q8U	1Q8W	1STC	1YDR	1YDS	1YDT	1JLU
y27	100	88	100	0	100	100	100	100
h52	100	100	100	0	100	100	100	8
m77	100	84	100	4	100	96	100	68
stu	100	100	100	100	100	100	100	100
iqp	4	0	92	100	100	88	52	24
iqs	32	0	68	48	88	96	80	56
iqb	72	84	44	4	44	60	88	24
avg.	72.57	65.14	86.29	36.57	90.29	91.43	88.57	54.29

Legend	[70% , 100%]
	[30%, 70%)
	[0%, 30%)
success rate	

the top-ranked structure has been used. Prior to docking, all seven ligands were minimized in vacuo using the MMFF94 force-field (Halgren 1996) to prevent the use of poor ligand geometries during docking. All results were averaged over 25 independent experiments and carried out on a single core of a Dual Core AMD Opteron 870, 2.0 GHz CPU using PLANTS standard settings and scoring function CHEMPLP (scoring of weak CH-O interactions activated) with $\sigma = 1.0$ for rigid receptor docking and $\sigma = 1.5$ for flexible receptor docking to account for the additional degrees of freedom.

Table 2 presents the results for docking the ligands into rigid (upper part) and flexible protein structures (lower part). For the docking into rigid protein structures, it can be ob-

served that all ligands can be docked back correctly into their native protein structures (bold entries) with high success rates. As already expected, ligand staurosporine (*stu*) is docked correctly into its native protein structure (PDB entry 1STC), but it is docked incorrectly in all others. On average, the protein structures of PDB-entries 1JLU (*apo form* of PKA), 1Q8T, 1Q8U, and 1STC perform worst with respect to their ability to reproduce the ligand poses correctly with success rates between 41% and 56%. In contrast, the protein structures of 1Q8W, 1YDR, 1YDS, and 1YDT are able to reproduce the correct conformations for all ligands except ligand staurosporine with average success rates of about 74% to 80%.

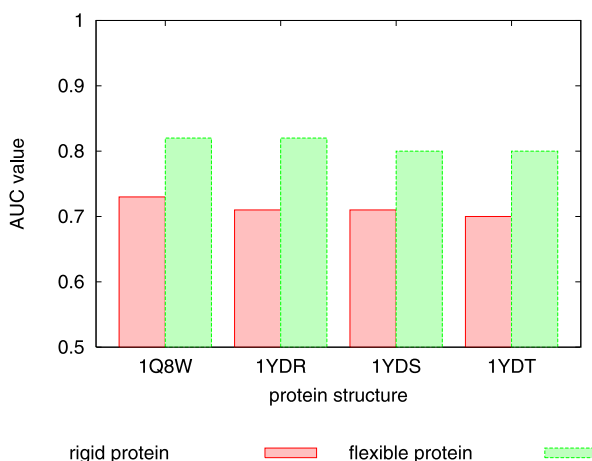
In the case of side-chain flexibility, the poses of all ligands can still be reproduced correctly in their native protein structures (bold entries). Only the success rate for docking ligand *iqs* back into its native protein structure (PDB-entries 1YDS) is insignificantly lowered when compared to the docking into the rigid protein structure. Ligand staurosporine clearly benefits from the protein flexibility as the correct pose can be reproduced for all protein structures now. In fact, flexibility leads to generally higher average success rates for each protein structure, except for the protein structure of PDB-entry 1STC. However, for various combinations of proteins and ligands in this experiment, the success rates observed are slightly worse than in the case when the protein is held rigid. (This lower success rate can be explained by the extended complexity of the problem and by inaccuracies in the intra-protein scoring function, which has to score the interactions arising between the flexible side-chain and the rigid part of the protein correctly. In the flexible protein side-chain case, the position of this side-chain needs to be correctly identified by the algorithm—in the rigid case this side-chain is already in a reasonably good position with respect to the ligand, with the main exception being for ligand staurosporine, as explained above.) Despite this fact, it is a very noteworthy result that the treatment of protein side-chain flexibility is useful to allow for a better reproduction of the correct ligand poses. In fact, for four out of eight protein structures (PDB-entries 1Q8W, 1YDR, 1YDS, and 1YDT) we were able to reproduce the poses of ligands inside 2.5 Å of the superimposed crystal structures with average success rates ranging from about 86% to 91%.

6.2 Virtual screening

In a next step, the influence of protein flexibility has been investigated in the context of *virtual screening*. The *virtual screening* of large compound libraries is one of the main applications of current docking tools. Therefore, PLANTS was also tested with respect to its ability to discriminate between biologically active and inactive ligands. The best-performing protein structures identified in the *cross-docking* of *protein kinase A* (PDB-entries 1Q8W, 1YDR, 1YDS, and 1YDT) have been used as the target structures. The compound library consisted of seven active (the ones already described in the PKA *cross-docking*) and 693 supposed inactive ligands. This active/inactive ratio was chosen, similar to many other virtual screening studies, to simulate the outcome of a real-world scenario, where usually less than 1% biologically active molecules are retrieved. Inactive ligands that approximately match the properties (number of rotatable bonds, hydrogen bond donors, acceptors and non-hydrogen atoms) of the active ligands were chosen from the ZINC database (Irwin and Shoichet 2005) to avoid artificial enrichment, that is, to avoid a bias in the results simply because clearly inappropriate ligands (for example, too small ones) were used as potential candidates (Verdonk et al. 2004). Like the seven ligand structures in the *cross-docking*, also the 693 inactive compounds were minimized in vacuo using the MMFF94 force-field (Halgren 1996) prior to docking.

The complete library of 700 ligands was then docked with PLANTS using standard settings and scoring function CHEMPLP (rigid receptor: $\sigma = 1$; partially flexible receptor:

Fig. 6 Comparison of AUC values obtained for the *virtual screening* against four *protein kinase A* (PKA) structures (PDB-entries 1Q8W, 1YDR, 1YDS, and 1YDT) considering rigid and partially flexible protein structures



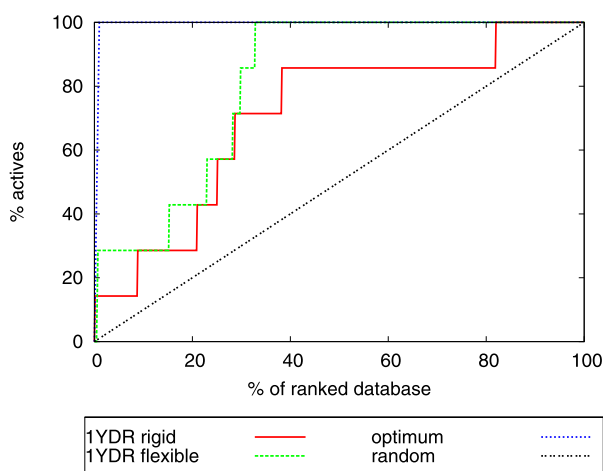
$\sigma = 1.5$) on a single core of a Dual Core AMD Opteron 870, 2.0 GHz CPU. The average docking times per ligand reached from about 27 s when considering rigid protein structures to approximately 65 s when considering partially flexible protein structures. After screening, the docked structures were sorted according to their scoring function value resulting in a ranked database. In a real-world screening, a fixed percentage representing the top-scoring ligands of the database would be used for *in vitro* tests for biological activity. Thus, for an algorithm to be useful in *virtual screening*, it is important that a percentage of active ligands that is as high as possible is within a small percentage of the top-ranked ligands the algorithm determines.

A common measure for evaluating the outcome of a *virtual screening* experiment is the *area under curve* (AUC) value, which is derived from the *receiver operating characteristic* (ROC) (Triballeau et al. 2005). Let TP (respectively, FP) be the *true (false) positives*, that is, the number of ligands selected by the docking tool that show (do not show) biological activity in *in vitro* tests, and FN (resp. TN) be the *false (true) negatives*, that is, ligands that are discarded by the docking protocol but show (do not show) biological activity. Then, the AUC value corresponds to the area under curve when plotting the *false positive rate* $1 - sp$ along the x -axis, where $sp = \frac{TN}{FP+TN}$ is also called the *specificity*, versus the *true positive rate se* along the y -axis, where $se = \frac{TP}{TP+FN}$ is the *sensitivity*. The maximum AUC value reachable is one and it corresponds to a perfect discrimination between active and inactive compounds, while AUC values around 0.5 show the performance of random selection.

The AUC values for all eight *virtual screening* campaigns are presented in Fig. 6. All four protein structures perform reasonably well in the case of docking into rigid protein structures with AUC values between 0.7 (PDB-entry 1YDT) and 0.73 (PDB-entry 1Q8W). An even better discrimination can be achieved when treating the side-chain of residue PHE327 flexible in the protein structure as already observed for the *cross-docking* experiments. In this case, the protein structures of 1Q8W and 1YDR perform best with an AUC value of 0.82. However, also for the other protein structures large improvements for the AUC value can be observed.

The enrichment curves for the protein structure of 1YDR comparing the rigid and the flexible case are shown in Fig. 7. Both docking protocols perform clearly better than random selection, but clearly make no optimal choices. For the rigid case, labeled as *1YDR rigid*, the biologically active ligands can be identified till up to 82% of the ranked database, which is

Fig. 7 Enrichments for the virtual screening against protein kinase A (PKA), PDB-entry 1YDR, considering rigid and partially flexible protein structures



apparently suboptimal. In contrast, the introduction of protein flexibility (labeled as *1YDR flexible*) allows for the identification of all biologically active ligands till about 33% of the ranked database, which is also reflected in the higher AUC value of 0.82.

7 Conclusions

We presented an effective docking algorithm based on the ACO metaheuristic. While most ACO algorithms are designed to tackle discrete optimization problems, we showed that such an approach, combined with a local search algorithm that works in the continuous search space, can be very effective to minimize real-valued objective functions such as they occur in the protein–ligand docking problem (PLDP). These results also suggest that algorithms similar to the proposed one can be effectively applied in other problem domains dealing with continuous function optimization. The experimental results have shown that PLANTS is a very effective tool for tackling the PLDP. The default settings of PLANTS are able to reproduce ligand geometries similar to the crystal geometry in about 72% of the cases for the structures of the CCDC/Astex benchmark set at docking times that are reasonable for practical applications. A direct comparison of PLANTS with the state-of-the-art docking program GOLD, which is based on a genetic algorithm, showed that PLANTS reaches higher pose prediction accuracy at similar or even lower docking times. While most current-generation docking algorithms treat only the ligand flexible, PLANTS additionally allows for partial protein flexibility by searching the torsional space of selected protein side-chains. The *cross-docking* and *virtual screening* experiments on target *protein kinase A* confirmed that the inclusion of additional degrees of freedom in docking calculations, like flexible side-chains in this case, can improve pose prediction success rates and enrichment rates significantly. Hence, the computational results confirm that Ant Colony Optimization algorithms can contribute towards the state-of-the-art for solving the practically relevant protein–ligand docking problem.

There are a number of directions for further research. Of significant scientific interest is certainly a detailed experimental study of the impact of various components of high-performance docking tools. Most noteworthy are here the scoring function chosen and the

search algorithm chosen. However, the successful docking tools use a number of problem-specific techniques that sometimes are not directly transferable when moving from one scoring function to another or from one search algorithm to another. Hence, such a study is well beyond the scope of this article and we intend to tackle this question in future research efforts. Another promising line of research is the further improvement of scoring functions. Even the best docking algorithms sometimes completely fail on specific target classes, which can in most cases be attributed to the scoring function. In this context, computationally more demanding scoring function terms like interactions between ligands and explicit water molecules are currently under investigation.

An executable version of PLANTS is available on request by contacting the authors.

Acknowledgements The authors thank Jens Gimmler and Nicola Zonta for helpful discussions and a careful reading of the manuscript; they also thank the three anonymous referees and the editor-in-chief for the comments that helped to improve the first version of the article. This work was supported by a scholarship of the Landesgraduiertenförderung Baden-Württemberg awarded to Oliver Korb. Thomas Stützle acknowledges support of the Belgian FNRS, of which he is a research associate.

References

- Abagyan, R., Totrov, M., & Kuznetsov, D. (1994). ICM—A new method for protein modeling and design: applications to docking and structure prediction from the distorted native conformation. *Journal of Computational Chemistry*, 15(5), 488–506.
- Baxter, C. A., Murray, C. W., Clark, D. E., Westhead, D. R., & Eldridge, M. D. (1997). Flexible docking using tabu search and an empirical estimate of binding affinity. *Proteins*, 33, 367–382.
- Berman, H., Westbrook, J., Feng, Z., Gilliland, G., Bhat, T., Weissig, H., et al. (2000). The protein data bank. *Nucleic Acids Research*, 28, 235–242.
- Clark, M., Cramer III, R. D., & van Opdenbosch, N. (1989). Validation of the general purpose tripos 5.2 force field. *Journal of Computational Chemistry*, 10, 982–1012.
- Dorigo, M., & Stützle, T. (2004). *Ant colony optimization*. Cambridge: MIT Press.
- Fischer, E. (1894). Einfluss der Configuration auf die Wirkung der Enzyme. *Chemische Berichte*, 27, 2985–2993.
- Gehlhaar, D. K., Verkhivker, G. M., Rejto, P. A., Sherman, C. J., Fogel, D. B., Fogel, L. J., & Freer, S. T. (1995). Molecular recognition of the inhibitor AG-1243 by HIV-1 protease: conformationally flexible docking by evolutionary programming. *Chemistry and Biology*, 2, 317–324.
- Halgren, T. A. (1996). Merck molecular force field, I: basis, form, scope, parameterization, and performance of MMFF94. *Journal of Computational Chemistry*, 17(5–6), 490–519.
- Irwin, J. J., & Shoichet, B. K. (2005). ZINC—A free database of commercially available compounds for virtual screening. *Journal of Chemical Information and Modeling*, 45(1), 177–82.
- Jones, G., Willett, P., Glen, R. C., Leach, A. R., & Taylor, R. D. (1997). Development and validation of a genetic algorithm for flexible docking. *Journal of Molecular Biology*, 267, 727–748.
- Kellenberger, E., Rodrigo, J., Muller, P., & Rognan, D. (2004). Comparative evaluation of eight docking tools for docking and virtual screening accuracy. *Proteins*, 57(2), 225–242.
- Kontoyianni, M., McClellan, L. M., & Sokol, G. S. (2004). Evaluation of docking performance: comparative data on docking algorithms. *Journal of Medicinal Chemistry*, 47(3), 558–565.
- Koshland, D. E. (1973). Protein shape and biological control. *Scientific American*, 229(4), 52–64.
- Korb, O., Stützle, T., & Exner, T. E. (2006). PLANTS: Application of Ant Colony Optimization to Structure-Based Drug Design. In M. Dorigo, L. M. Gambardella, M. Birattari, A. Martinoli, R. Poli, & T. Stützle (Eds.), *Lecture Notes in Computer Science: Vol. 4150. Ant colony optimization and swarm intelligence, 5th international workshop, ANTS 2006* (pp. 247–258). Berlin: Springer.
- McMartin, C., & Bohacek, R. S. (1997). QXP: powerful, rapid computer algorithms for structure-based drug design. *Journal of Computer-Aided Molecular Design*, 11(4), 333–344.
- Morris, G. M., Goodsell, D. S., Halliday, R. S., Huey, R., Hart, W. E., Belew, R. K., et al. (1998). Automated docking using a Lamarckian genetic algorithm and an empirical binding free energy function. *Journal of Computational Chemistry*, 19, 1639–1662.
- Müller, G. (2003). Medicinal chemistry of target family-directed masterkeys. *Drug Discovery Today*, 8(15), 681–691.

- Nelder, J. A., & Mead, R. (1965). A simplex method for function minimization. *Computer Journal*, 7, 308–313.
- Nissink, J. W. M., Murray, C., Hartshorn, M., Verdonk, M. L., Cole, J. C., & Taylor, R. (2002). A new test set for validating predictions of protein–ligand interaction. *Proteins*, 49(4), 457–471.
- Press, W. H., Flannery, B. P., Teukolsky, S. A., & Vetterling, W. T. (1992). *Numerical recipes in C: the art of scientific computing*. Cambridge: Cambridge University Press.
- Stützle, T., & Hoos, H. H. (2000). $\mathcal{M}\mathcal{A}\mathcal{X}$ – $\mathcal{M}\mathcal{I}\mathcal{N}$ ant system. *Future Generation Computer Systems*, 16(8), 889–914.
- Taylor, R. D., Jewsbury, P. J., & Essex, J. W. (2002). A review of protein–small molecule docking methods. *Journal of Computer-Aided Molecular Design*, 16, 151–166.
- Triballeau, N., Acher, F., Brabet, I., Pin, J.-P., & Bertrand, H.-O. (2005). Virtual screening workflow development guided by the “receiver operating characteristic” curve approach. Application to high-throughput docking on metabotropic glutamate receptor subtype 4. *Journal of Medicinal Chemistry*, 48(7), 2534–2547.
- Verdonk, M. L., Cole, J. C., Hartshorn, M. J., Murray, C. W., & Taylor, R. D. (2003). Improved protein–ligand docking using GOLD. *Proteins*, 52, 609–623.
- Verdonk, M. L., Berdini, V., Hartshorn, M. J., Mooij, W. T. M., Murray, C. W., Taylor, R. D., & Watson, P. (2004). Virtual screening using protein–ligand docking: avoiding artificial enrichment. *Journal of Chemical Information and Modeling*, 44(3), 793–806.
- Verkhivker, G. M. (2004). Computational analysis of ligand binding dynamics at the intermolecular hot spots with the aid of simulated tempering and binding free energy calculations. *Journal of Molecular Graphics and Modelling*, 22, 335–348.